



SLLIM: System Log Local Intelligent Model

Authors: Carlos Cruzportillo, Nassos Galiopoulos, Jason Gillette

Affiliation: University of Texas at San Antonio

Date: December 2nd, 2024

Introduction

Problem Statement

The increasing volume of system logs generated by interconnected devices and enterprise systems creates a challenge for IT professionals in efficiently detecting threats and diagnosing issues, necessitating the development of lightweight, intelligent tools for real-time log analysis and query.

Do more with less!

Introduction

Specific Objectives

1. Fine-tune at least two lightweight LLMs for comparative analysis. ❌
2. Evaluate question answering performance of lightweight LLMs versus resource-intensive models in the cybersecurity domain. ✅

Introduction

Research Questions

1. How well can lightweight LLMs detect system issues and security threats from system logs? ✓
2. How effectively can lightweight LLMs perform question answering compared to larger, more resource-intensive models? ✓

Methods

Models

- **meta-llama/Llama-3.1-8B** as a **Large** Subject
- **meta-llama/Llama-3.2-1B** as a *small* Subject

Access to both models was granted through *Meta AI* via *HuggingFace* 😊

Model

Llama-3.1-8B

- Multi-lingual text LLM
- Released July 2024
- 8 billion parameters
- ~32 GB in size at FP32 (*4 bytes per param*)

Model

Llama-3.2-1B

- Multi-lingual text LLM
- Released September 2024
- 1 billion parameters
- ~4 GB in size at FP32 (*4 bytes per param*)

Dataset

- **Source:** LogQA, a question-answering dataset derived from three public log datasets: HDFS, OpenSSH, and Spark.
- **Composition:**
 - **Raw Logs:** 2,000 log entries selected per dataset.
 - **Question Generation:** Utilized a question generation model to create reading comprehension-style questions with answers extracted from the logs.
- **Example:**

```
{  
  "Question": "What is the status of the block blk_-6369730481066968769?",  
  "Answer": "terminating",  
  "RawLog": "PacketResponder 1 for block blk_-6369730481066968769 terminating"  
}
```


Methods

Evaluation Framework

1. Evaluation Metrics

- **Exact Match (EM)**: Binary match between the generated answer and ground truth.
- **Contains Match (CM)**: Evaluates if the generated answer contains the ground truth.
- **Token-based F1**: Precision and recall on aligned tokens.
- **BERTScore**: Semantic similarity using embeddings.

Methods

Experiment Setup

Few-Shot vs Zero-Shot

- **Few-Shot Prompting:**
 - Incorporates 2 examples from the training set as context.
 - Example: "*Context: [Log 1] Question: [Q1] Answer: [A1]*".
- **Zero-Shot Prompting:**
 - Provides no additional examples.
 - Question directly follows the context.

Methods

Inference Workflow

1. Model Loading:

- HuggingFace pipelines with `transformers`.
- Quantized models using `BitsAndBytesConfig` for memory optimization.

2. Batch Processing:

- Test data processed in batches of 2 for GPU efficiency.
- Token truncation managed dynamically (`max_new_tokens`).

3. Data Storage:

- Results saved as structured JSON for easy evaluation.
- Includes metadata: context, question, generated answer.

Why GPU efficiency matters

NVIDIA-SMI 565.72			Driver Version: 566.14			CUDA Version: 12.7		
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.	MIG M.
0	NVIDIA GeForce RTX 2060		On	00000000:01:00.0	On			N/A
N/A	<u>74C</u>	P0	69W / 80W	5849MiB / 6144MiB		<u>100%</u>	Default	N/A

Why GPU efficiency matters

NVIDIA-SMI 565.72			Driver Version: 566.14			CUDA Version: 12.7		
GPU	Name		Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr. ECC
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage		GPU-Util	Compute M. MIG M.
0	NVIDIA GeForce RTX 2060		On		00000000:01:00.0	On		N/A
N/A	73C	P0	73W / 80W		1334MiB / 6144MiB		41%	Default
							<div><div></div></div>	N/A

Results

Aggregated Metrics

Metric	Llama-3.1- 8B ZS	Llama-3.2- 1B ZS	Llama-3.1- 8B FS	Llama-3.2- 1B FS
Exact Match	0.00	0.02	0.01	0.02
Contains Match	0.71	0.58	0.51	0.62
Token F1	0.16	0.22	0.16	0.20
BERTScore	0.83	0.82	0.77	0.79

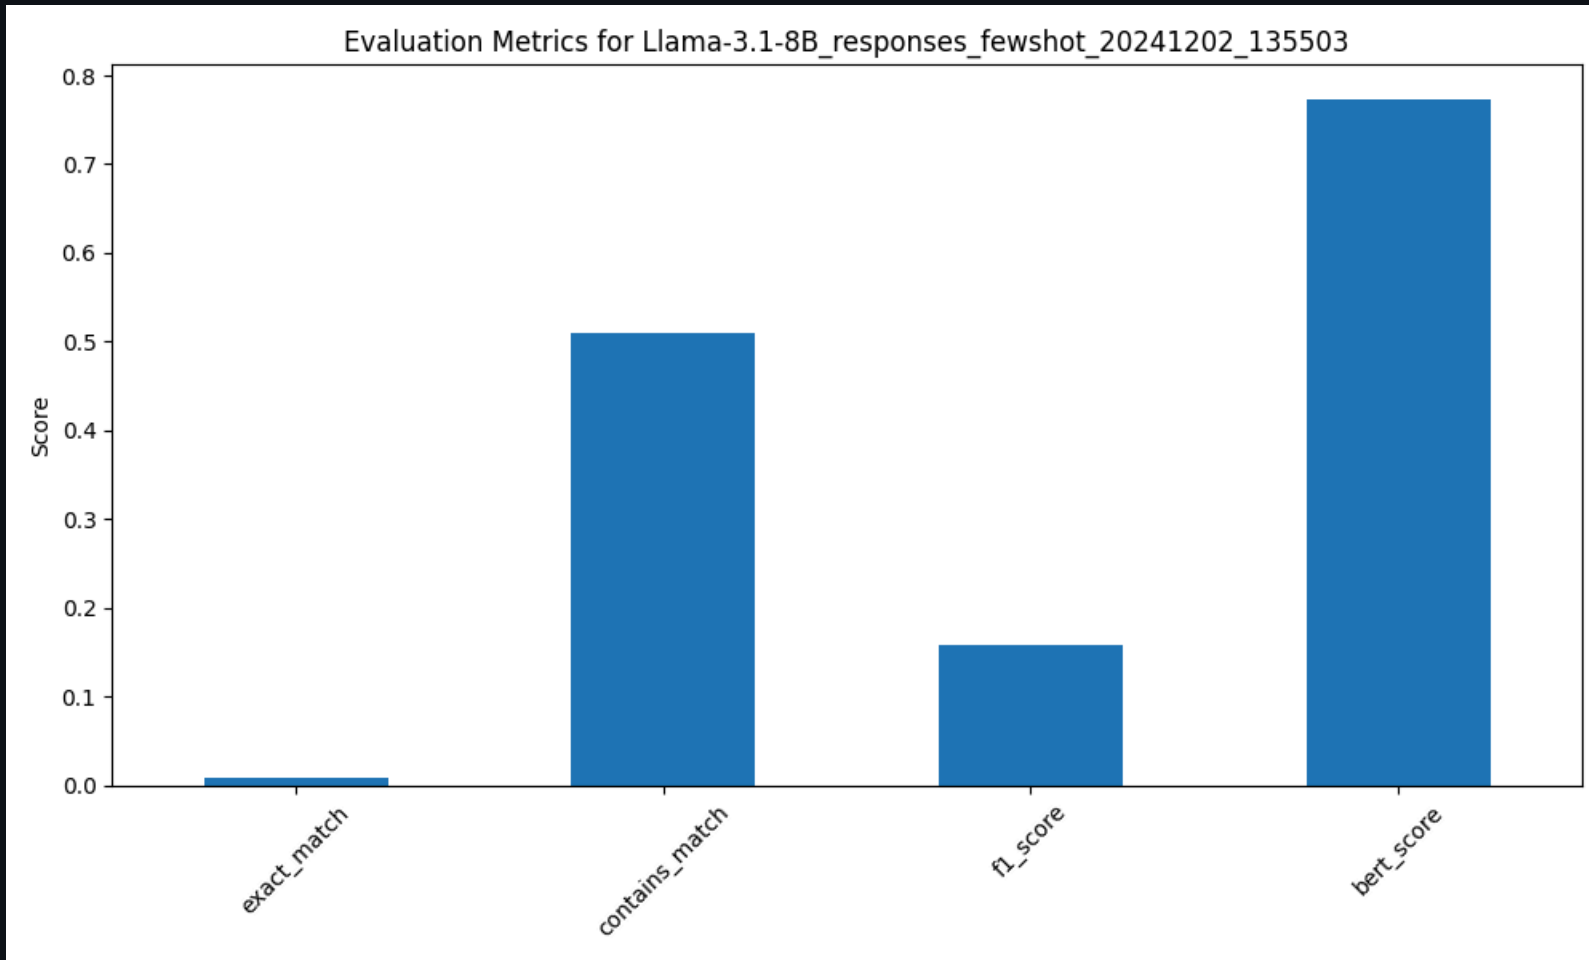
Results (cont.)

Observations

- Larger models exhibited only marginally better performance.
- Few-shot prompting degraded performance.
- Tokenization likely skewed matches and underrepresented performance.

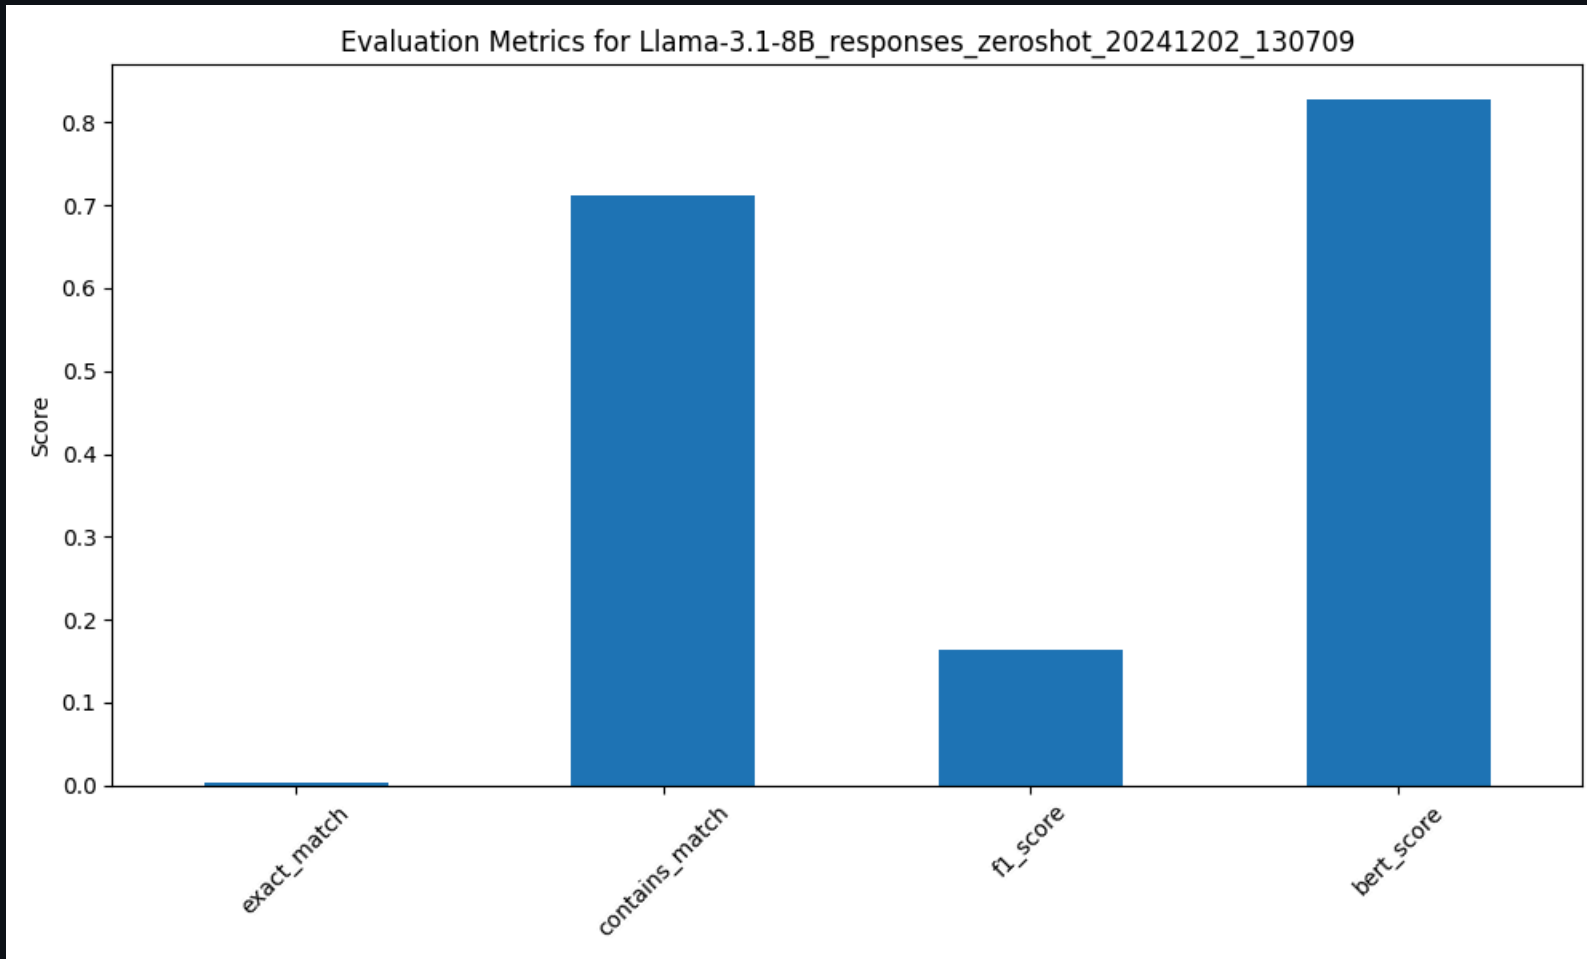
Results (cont.)

Metric Comparison




Results (cont.)

Metric Comparison




Results (cont.)

Metric Comparison



Results (cont.)

Metric Comparison



Key Findings

1. Model Efficiency:

- Llama-3.2-1B successfully handles domain-specific tasks with minimal resources.
- Llama-3.1-8B demonstrates no significant performance gain.

2. Few-Shot Effectiveness:

- Few-shot prompting introduces noise and complexity, for no performance gain.

3. Challenges:

- Token alignment in noisy log contexts affects EM and F1 scores.
- BERTScore highlights semantic drift in generated answers.
- BERTScore does not tell the full story as wrong answers can be

Conclusion

Contributions

1. A comparative analysis of **lightweight** and **large-scale** LLMs in log-based QA tasks.
2. A novel application of lightweight LLMs on Log analysis tasks.

Next Steps

1. Fine-tuning vs. in-context learning.
2. Improving / validating evaluation.
3. Exploring other evaluation metrics for log anomaly detection.
4. Expand scope of log analysis to threat detection.

Questions?